

LIN7076 – Foundations of Computational Linguistics

Linear Regression

Adèle Hénot-Mortier

14/10/2025

Queen Mary University of London

Just Five Straight Minutes Of Obama Pauses

What do you think are factors (linguistic, sociological, contextual) influencing how long you pause after you end a sentence?

How to **predict** a continuous variable based on **multiple factors**?

- For instance, how to predict the duration of a **pause** at the end of a sentence in a speech, given:
 - the **length** of the sentence (number of words);
 - the **speech rate** of the speaker (in words/minute);
 - the **context** of the speech (formal/informal);
 - the **gender** of the speaker (F/M/other);
 - the **importance** of the topic (unimportant/moderately important/very important).

Plan for today

- Define the **kinds of variables** involved in a linear regression model.
- Define what it means for a function to be **linear**.
- Spell-out how non-numerical variables can be **encoded** and used in a linear model.
- Discuss how a linear model can be **used and interpreted**.
- Explain how a linear model can be **fitted** to the data (time allowing).
- Discuss an improvement of the basic model, called **regularization** (time allowing).

Defining the variables of a linear model

Two main kinds of variable

- The variable we want to predict (e.g. pause duration) is called a **dependent** variable.
- The variables we use to predict it (e.g. sentence length, speech rate, context, gender, topic importance) are called **independent** variables.

A linear regression predicts the **dependent** variable from the **independent** variable(s) assuming their relation is **linear**.

$$\widehat{\text{pause}} = f(\text{length}, \text{speech_rate}, \text{context}, \text{gender}, \text{importance})$$

With f a “linear” function.

Two main kinds of variable

- The variable we want to predict (e.g. pause duration) is called a **dependent** variable.
- The variables we use to predict it (e.g. sentence length, speech rate, context, gender, topic importance) are called **independent** variables.

A linear regression predicts the **dependent** variable from the **independent** variable(s) assuming their relation is **linear**.

$$\widehat{\text{pause}} = f(\text{length}, \text{speech_rate}, \text{context}, \text{gender}, \text{importance})$$

With f a “linear” function.

Two main kinds of variable

- The variable we want to predict (e.g. pause duration) is called a **dependent** variable.
- The variables we use to predict it (e.g. sentence length, speech rate, context, gender, topic importance) are called **independent** variables.

A linear regression predicts the **dependent** variable from the **independent** variable(s) assuming their relation is **linear**.

$$\widehat{\text{pause}} = f(\text{length}, \text{speech_rate}, \text{context}, \text{gender}, \text{importance})$$

With f a “linear” function.

Types of independent variables

- Different independent variables are encoded differently!
- **Quantitative variables** represent “real” amounts that can be combined *via* arithmetic operations (+, - etc.). Two subtypes:
 - Discrete, integer-like; e.g. number of words (1, 2, 3, 4...).
 - Continuous, e.g. speech rate, pause duration (0, ..., 1.5, 1.6, 1.7, ...).
- **Categorical variables** represent groupings. Three subtypes:
 - Binary variables with 2 levels, e.g. a formal vs. informal context.
 - Nominal, unordered variables with > 2 levels, e.g. genders.
 - Ordinal, ordered variables with > 2 levels, e.g. unimportant, moderately important, and very important topic. The distance between each level may be variable and subjective.

What about... positive/negative sentiment? age in years?
socio-economic status (> 1 possibility)? Part-of-Speech (Noun, Verb, Adjective...)?

Types of independent variables

- Different independent variables are encoded differently!
- **Quantitative variables** represent “real” amounts that can be combined *via* arithmetic operations (+, - etc.). Two subtypes:
 - **Discrete**, integer-like; e.g. number of words (1, 2, 3, 4...).
 - **Continuous**, e.g. speech rate, pause duration (0, ..., 1.5, 1.6, 1.7, ...).
- **Categorical variables** represent groupings. Three subtypes:
 - **Binary** variables with 2 levels, e.g. a formal vs. informal context.
 - **Nominal**, unordered variables with > 2 levels, e.g. genders.
 - **Ordinal**, ordered variables with > 2 levels, e.g. unimportant, moderately important, and very important topic. The distance between each level may be variable and subjective.

What about... positive/negative sentiment? age in years?
socio-economic status (> 1 possibility)? Part-of-Speech (Noun, Verb, Adjective...)?

Types of independent variables

- Different independent variables are encoded differently!
- **Quantitative variables** represent “real” amounts that can be combined *via* arithmetic operations (+, - etc.). Two subtypes:
 - **Discrete**, integer-like; e.g. number of words (1, 2, 3, 4...).
 - **Continuous**, e.g. speech rate, pause duration (0, ..., 1.5, 1.6, 1.7, ...).
- **Categorical variables** represent groupings. Three subtypes:
 - **Binary** variables with 2 levels, e.g. a formal vs. informal context.
 - **Nominal**, unordered variables with > 2 levels, e.g. genders.
 - **Ordinal**, ordered variables with > 2 levels, e.g. unimportant, moderately important, and very important topic. The distance between each level may be variable and subjective.

What about... positive/negative sentiment? age in years?
socio-economic status (> 1 possibility)? Part-of-Speech (Noun, Verb, Adjective...)?

Types of independent variables

- Different independent variables are encoded differently!
- **Quantitative variables** represent “real” amounts that can be combined *via* arithmetic operations (+, - etc.). Two subtypes:
 - **Discrete**, integer-like; e.g. number of words (1, 2, 3, 4...).
 - **Continuous**, e.g. speech rate, pause duration (0, ..., 1.5, 1.6, 1.7, ...).
- **Categorical variables** represent groupings. Three subtypes:
 - **Binary** variables with 2 levels, e.g. a formal vs. informal context.
 - **Nominal**, unordered variables with > 2 levels, e.g. genders.
 - **Ordinal**, ordered variables with > 2 levels, e.g. unimportant, moderately important, and very important topic. The distance between each level may be variable and subjective.

What about... positive/negative sentiment? age in years?
socio-economic status (> 1 possibility)? Part-of-Speech (Noun, Verb, Adjective...)?


Defining linear relations between variables

Linear relations

A linear regression predicts the dependent variable from the independent variable(s) assuming their relation is **linear**.


- Two variables x and y are linearly related if changes in x are directly proportional to changes in y at a constant rate β .

$$y = \beta x + \alpha$$




Made with ChatGPT

Linear relation practice 1




Linear relation practice 1




- Curve crosses y -axis at $y = +1$, so intercept α is $+1$;
- Curve goes $+3$ on y for each $+1$ on x : slope β is thus $3/1 = 3$. It's **positive**, consistent with the fact that the curve is **increasing**.
- So the linear function represented here is $y = 3x + 1$.

Linear relation practice 2




Linear relation practice 2




- Curve crosses y -axis at $y = -2$, so intercept α is -2 ;
- Curve goes $+1$ on y for each $+2$ on x : slope β is thus $1/2 = .5$. It's **positive**, consistent with the fact that the curve is **increasing**.
- So the linear function represented here is $y = .5x - 2$.

Linear relation practice 3



Linear relation practice 3




- Curve crosses y-axis at $y = +3$, so intercept α is $+3$;
- Curve goes -2 on y for each $+1$ on x : slope β is thus $-2/1 = -2$. It's **negative**, consistent with the fact that the curve is **decreasing**.
- So the linear function represented here is $y = -2x + 3$.

Approximating linear relations: negative correlation

- One can look at **how a single independent variable influences the dependent variable**, e.g. how speech rate (continuous) influences pause duration (continuous).

Pause duration as a function of speech rate (MADE UP DATA)




$$\underbrace{\hat{y}}_{\text{pause}} = \beta \underbrace{x}_{\text{rate}} + \alpha$$

Approximating linear relations: negative correlation

- One can look at **how a single independent variable influences the dependent variable**, e.g. how speech rate (continuous) influences pause duration (continuous).

Pause duration as a function of speech rate (MADE UP DATA)




$$\underbrace{\hat{y}}_{\text{pause}} = \beta \underbrace{x}_{\text{rate}} + \alpha$$

- Sign of β ?

Approximating linear relations: negative correlation

- One can look at how a single independent variable influences the dependent variable, e.g. how speech rate (continuous) influences pause duration (continuous).




$$\underbrace{\hat{y}}_{\text{pause}} = \beta \underbrace{x}_{\text{rate}} + \alpha$$

- Sign of β ?
- We say that pause duration and speech rate are **negatively correlated**.

Approximating linear relations: positive correlation

- One can also look at how sentence length (quantitative, discrete) influences pause duration (quantitative, continuous).

Pause duration as a function of sentence length (MADE UP DATA)




$$\underbrace{\hat{y}}_{\text{pause}} = \beta \underbrace{x}_{\text{length}} + \alpha$$

Approximating linear relations: positive correlation

- One can also look at how sentence length (quantitative, discrete) influences pause duration (quantitative, continuous).

Pause duration as a function of sentence length (MADE UP DATA)



$$\underbrace{\hat{y}}_{\text{pause}} = \beta \underbrace{x}_{\text{length}} + \alpha$$

- Sign of β ?

Approximating linear relations: positive correlation

- One can also look at how sentence length (quantitative, discrete) influences pause duration (quantitative, continuous).

Pause duration as a function of sentence length (MADE UP DATA)



$$\underbrace{\hat{y}}_{\text{pause}} = \beta \underbrace{x}_{\text{length}} + \alpha$$

- Sign of β ?
- We say that pause duration and speech rate are **positively correlated**.

Building up the linear regression formula

- We “superpose” several one-to-one linear relations in one single formula. This is still a **linear relation**.

$$\hat{y} = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha$$

- y is a continuous, dependent variable, like pause duration.
- The x_i s are the independent variables, and the β_i s their respective slopes (**parameters** of the linear model).
 - If variable x_i **positively** influences y , then β_i will be **positive**.
 - If variable x_i **negatively** influences y , then β_i will be **negative**.
 - The **magnitude** of β_i ($|\beta_i|$) encodes the **strength** of the association.
- α is the intercept of the model (also a **parameter**). It can be seen as the **baseline value** for y , that would be observed if none of the independent variables were relevant, i.e. if all β_i were zero.

Building up the linear regression formula

- We “superpose” several one-to-one linear relations in one single formula. This is still a **linear relation**.

$$\hat{y} = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha$$

- y is a continuous, dependent variable, like pause duration.
- The x_i s are the independent variables, and the β_i s their respective slopes (**parameters** of the linear model).
 - If variable x_i **positively** influences y , then β_i will be **positive**.
 - If variable x_i **negatively** influences y , then β_i will be **negative**.
 - The **magnitude** of β_i ($|\beta_i|$) encodes the **strength** of the association.
- α is the intercept of the model (also a **parameter**). It can be seen as the **baseline value** for y , that would be observed if none of the independent variables were relevant, i.e. if all β_i were zero.

Building up the linear regression formula

- We “superpose” several one-to-one linear relations in one single formula. This is still a **linear relation**.

$$\hat{y} = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha$$

- y is a continuous, dependent variable, like pause duration.
- The x_i s are the independent variables, and the β_i s their respective slopes (**parameters** of the linear model).
 - If variable x_i **positively** influences y , then β_i will be **positive**.
 - If variable x_i **negatively** influences y , then β_i will be **negative**.
 - The **magnitude** of β_i ($|\beta_i|$) encodes the **strength** of the association.
- α is the intercept of the model (also a **parameter**). It can be seen as the **baseline value** for y , that would be observed if none of the independent variables were relevant, i.e. if all β_i were zero.

Building up the linear regression formula

- We “superpose” several one-to-one linear relations in one single formula. This is still a **linear relation**.

$$\hat{y} = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha$$

- y is a continuous, dependent variable, like pause duration.
- The x_i s are the independent variables, and the β_i s their respective slopes (**parameters** of the linear model).
 - If variable x_i **positively** influences y , then β_i will be **positive**.
 - If variable x_i **negatively** influences y , then β_i will be **negative**.
 - The magnitude of β_i ($|\beta_i|$) encodes the **strength** of the association.
- α is the intercept of the model (also a **parameter**). It can be seen as the **baseline value** for y , that would be observed if none of the independent variables were relevant, i.e. if all β_i were zero.

Building up the linear regression formula

- We “superpose” several one-to-one linear relations in one single formula. This is still a **linear relation**.

$$\hat{y} = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha$$

- y is a continuous, dependent variable, like pause duration.
- The x_i s are the independent variables, and the β_i s their respective slopes (**parameters** of the linear model).
 - If variable x_i **positively** influences y , then β_i will be **positive**.
 - If variable x_i **negatively** influences y , then β_i will be **negative**.
 - The magnitude of β_i ($|\beta_i|$) encodes the **strength** of the association.
- α is the intercept of the model (also a **parameter**). It can be seen as the **baseline value** for y , that would be observed if none of the independent variables were relevant, i.e. if all β_i were zero.

Building up the linear regression formula

- We “superpose” several one-to-one linear relations in one single formula. This is still a **linear relation**.

$$\hat{y} = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha$$

- y is a continuous, dependent variable, like pause duration.
- The x_i s are the independent variables, and the β_i s their respective slopes (**parameters** of the linear model).
 - If variable x_i **positively** influences y , then β_i will be **positive**.
 - If variable x_i **negatively** influences y , then β_i will be **negative**.
 - The **magnitude** of β_i ($|\beta_i|$) encodes the **strength** of the association.
- α is the intercept of the model (also a **parameter**). It can be seen as the **baseline value** for y , that would be observed if none of the independent variables were relevant, i.e. if all β_i were zero.

Sketching a linear model of pause duration

$$\widehat{\text{pause}} = \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} + \beta_4 \text{gender} \\ + \beta_5 \text{importance} + \alpha$$

Sketching a linear model of pause duration

$$\widehat{\text{pause}} = \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} + \beta_4 \text{gender} \\ + \beta_5 \text{importance} + \alpha$$

- Suppose we have a sentence of **length** 15, and the speaker has a speech **rate** of 120 words per minute. We have $\beta_1 = 2/3$, $\beta_2 = -1/10$, $\alpha = 70$ (milliseconds), and all the other β s are 0. What's the estimated **pause** (in milliseconds)?

Sketching a linear model of pause duration

$$\widehat{\text{pause}} = \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} + \beta_4 \text{gender} \\ + \beta_5 \text{importance} + \alpha$$

- Suppose we have a sentence of **length** 15, and the speaker has a speech **rate** of 120 words per minute. We have $\beta_1 = 2/3$, $\beta_2 = -1/10$, $\alpha = 70$ (milliseconds), and all the other β s are 0. What's the estimated **pause** (in milliseconds)?

$$\widehat{\text{pause}} = 2/3 \times 15 + 1/10 \times 120 + 70 = 10 - 12 + 70 = 68 \text{ milliseconds}$$

What still needs to be fleshed out

$$\widehat{\text{pause}} = \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} + \beta_4 \text{gender} \\ + \beta_5 \text{importance} + \alpha$$

- So far we've seen how to model dependencies between quantitative variables (numbers!), like $\widehat{\text{pause}} = f(\text{length})$, and $\widehat{\text{pause}} = f(\text{rate})$.
- What to do with categorical variables which represent groupings, like **context** (assumed binary), **gender** (assumed ternary/nominal), and **importance** (assumed ordinal)?

What still needs to be fleshed out

$$\widehat{\text{pause}} = \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} + \beta_4 \text{gender} \\ + \beta_5 \text{importance} + \alpha$$

- So far we've seen how to model dependencies between quantitative variables (numbers!), like $\widehat{\text{pause}} = f(\text{length})$, and $\widehat{\text{pause}} = f(\text{rate})$.
- What to do with categorical variables which represent groupings, like **context** (assumed binary), **gender** (assumed ternary/nominal), and **importance** (assumed ordinal)?

What still needs to be fleshed out

$$\widehat{\text{pause}} = \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} + \beta_4 \text{gender} \\ + \beta_5 \text{importance} + \alpha$$

- So far we've seen how to model dependencies between quantitative variables (numbers!), like $\widehat{\text{pause}} = f(\text{length})$, and $\widehat{\text{pause}} = f(\text{rate})$.
- What to do with categorical variables which represent groupings, like **context** (assumed binary), **gender** (assumed ternary/nominal), and **importance** (assumed ordinal)?

Encoding categorical variables

Encoding binary variables

- Binary variables can be encoded as **two different integers**, e.g. 0 and 1. In that case class 0 is the baseline class, and β quantifies the effect of class 1. One could also pick -1 and 1.
- For instance, if we assume the **context** of speech can be informal (c_0) or formal (c_1), the “context” variable will be 0 if the context is informal and 1 if it’s formal.

$$\begin{aligned}\text{context} &= \begin{cases} 0 & \text{if context is informal} \\ 1 & \text{if context is formal} \end{cases} \\ &= \mathbb{1}_{\text{formal}}\end{aligned}$$

- $\mathbb{1}_{\text{formal}}$ is called the **indicator function** for formal contexts.

Encoding binary variables

- Binary variables can be encoded as **two different integers**, e.g. 0 and 1. In that case class 0 is the baseline class, and β quantifies the effect of class 1. One could also pick -1 and 1.
- For instance, if we assume the **context** of speech can be informal (c_0) or formal (c_1), the “context” variable will be 0 if the context is informal and 1 if it’s formal.

$$\begin{aligned}\text{context} &= \begin{cases} 0 & \text{if context is informal} \\ 1 & \text{if context is formal} \end{cases} \\ &= \mathbb{1}_{\text{formal}}\end{aligned}$$

- $\mathbb{1}_{\text{formal}}$ is called the **indicator function** for formal contexts.

Encoding ordinal variables

- Ordinal variables like topic **importance**, can be **mapped to integers**. This is treating them as discrete.
- For instance, a unimportant/moderately important/very important scale can be encoded as 0/1/2.



- This encoding is assuming that the **“distance” between the different levels is uniform** (because the distance between 0 and 1, and 1 and 2, is the same).
- This is often good enough but if not, we go for a different encoding called **thermometer encoding** (see last slide of this presentation for more).

Encoding ordinal variables

- Ordinal variables like topic **importance**, can be **mapped to integers**. This is treating them as discrete.
- For instance, a unimportant/moderately important/very important scale can be encoded as 0/1/2.



- This encoding is assuming that the “distance” between the different levels is uniform (because the distance between 0 and 1, and 1 and 2, is the same).
- This is often good enough but if not, we go for a different encoding called **thermometer encoding** (see last slide of this presentation for more).

Encoding ordinal variables


- Ordinal variables like topic **importance**, can be **mapped to integers**. This is treating them as discrete.
- For instance, a unimportant/moderately important/very important scale can be encoded as 0/1/2.



- This encoding is assuming that **the “distance” between the different levels is uniform** (because the distance between 0 and 1, and 1 and 2, is the same).
- This is often good enough but if not, we go for a different encoding called **thermometer encoding** (see last slide of this presentation for more).

Encoding ordinal variables



- Ordinal variables like topic **importance**, can be **mapped to integers**. This is treating them as discrete.
- For instance, a unimportant/moderately important/very important scale can be encoded as 0/1/2.



- This encoding is assuming that **the “distance” between the different levels is uniform** (because the distance between 0 and 1, and 1 and 2, is the same).
- This is often good enough but if not, we go for a different encoding called **thermometer encoding** (see last slide of this presentation for more).

Linear dependencies with nominal variables



- What about the effect of **gender** (categorical, nominal) on pause duration—if any?
- Genders are not ordered, and so gender cannot be “increased” to see its influence on pause duration!



- Still, can we see a gender effect in the above data if we compare one class (e.g. M) to the others?

Linear dependencies with nominal variables



- What about the effect of **gender** (categorical, nominal) on pause duration—if any?
- Genders are not ordered, and so gender cannot be “increased” to see its influence on pause duration!



- Still, can we see a gender effect in the above data if we compare one class (e.g. M) to the others?



Linear dependencies with nominal variables

- What about the effect of **gender** (categorical, nominal) on pause duration—if any?
- Genders are not ordered, and so gender cannot be “increased” to see its influence on pause duration!





- Still, can we see a gender effect in the above data if we compare one class (e.g. M) to the others?

One vs. all comparisons with nominal variables




- These made-up data suggest that:
 - being a man leads to longer pauses than not being a man;
 - being a women leads to shorter pauses than not being a woman;
 - belonging to another gender is not very predictive due to a lot of intra-group variation.
- Just FYI apparently there is no significant effect of gender on pause/frequency duration. These data and made-up differences are just here for illustration.

One vs. all comparisons with nominal variables





- These made-up data suggest that:
 - being a man leads to longer pauses than not being a man;
 - being a women leads to shorter pauses than not being a woman;
 - belonging to another gender is not very predictive due to a lot of intra-group variation.
- Just FYI apparently there is no significant effect of gender on pause/frequency duration. These data and made-up differences are just here for illustration.

One vs. all comparisons with nominal variables




- These made-up data suggest that:
 - being a man leads to longer pauses than not being a man;
 - being a women leads to shorter pauses than not being a woman;
 - belonging to another gender is not very predictive due to a lot of intra-group variation.
- Just FYI apparently there is no significant effect of gender on pause/frequency duration. These data and made-up differences are just here for illustration.

One vs. all comparisons with nominal variables



- These made-up data suggest that:
 - being a man leads to longer pauses than not being a man;
 - being a women leads to shorter pauses than not being a woman;
 - belonging to another gender is not very predictive due to a lot of intra-group variation.
- Just FYI apparently there is no significant effect of gender on pause/frequency duration. These data and made-up differences are just here for illustration.

One vs. all comparisons with nominal variables



- These made-up data suggest that:
 - being a man leads to longer pauses than not being a man;
 - being a women leads to shorter pauses than not being a woman;
 - belonging to another gender is not very predictive due to a lot of intra-group variation.
- **Just FYI apparently there is no significant effect of gender on pause/frequency duration.** These data and made-up differences are just here for illustration.

Encoding nominal variables

- Assuming a M/F/Other tripartition for **gender**, we can encode this nominal variable using 2 binary variables:

Encoding nominal variables

- Assuming a M/F/Other tripartition for **gender**, we can encode this nominal variable using 2 binary variables:
 - $\text{is_male} = 1_M = 1$ iff the speaker is male;

Encoding nominal variables

- Assuming a M/F/Other tripartition for **gender**, we can encode this nominal variable using 2 binary variables:
 - $\text{is_male} = \mathbf{1}_M = 1$ iff the speaker is male;
 - and $\text{is_female} = \mathbf{1}_F = 1$ iff the speaker is female.

Encoding nominal variables

- Assuming a M/F/Other tripartition for **gender**, we can encode this nominal variable using 2 binary variables:
 - $\text{is_male} = \mathbb{1}_M = 1$ iff the speaker is male;
 - and $\text{is_female} = \mathbb{1}_F = 1$ iff the speaker is female.
- The last option is encoded by setting **is_male** and **is_female** to 0.

Encoding nominal variables

- Assuming a M/F/Other tripartition for **gender**, we can encode this nominal variable using 2 binary variables:
 - **is_male**= $\mathbb{1}_M=1$ iff the speaker is male;
 - and **is_female**= $\mathbb{1}_F=1$ iff the speaker is female.
- The last option is encoded by setting **is_male** and **is_female** to 0.
- More generally, a nominal variable with K levels ($c_0, \dots, c_{K-1}$), can be “split” into $K - 1$ **binary variables**: $\mathbb{1}_{c_0}$ (=1 iff the class is c_0), $\mathbb{1}_{c_1}$ (=1 iff the class is c_1), ..., $\mathbb{1}_{c_{K-2}}$ (=1 iff the class is c_{K-2}).

Encoding nominal variables

- Assuming a M/F/Other tripartition for **gender**, we can encode this nominal variable using 2 binary variables:
 - **is_male**= $\mathbb{1}_M=1$ iff the speaker is male;
 - and **is_female**= $\mathbb{1}_F=1$ iff the speaker is female.
- The last option is encoded by setting **is_male** and **is_female** to 0.
- More generally, a nominal variable with K levels ($c_0, \dots, c_{K-1}$), can be “split” into $K - 1$ **binary variables**: $\mathbb{1}_{c_0}$ (=1 iff the class is c_0), $\mathbb{1}_{c_1}$ (=1 iff the class is c_1), ..., $\mathbb{1}_{c_{K-2}}$ (=1 iff the class is c_{K-2}).
- Why not include $\mathbb{1}_{c_{K-1}}$, the indicator function for the last class?

Encoding nominal variables

- Assuming a M/F/Other tripartition for **gender**, we can encode this nominal variable using 2 binary variables:
 - $\text{is_male} = \mathbb{1}_M = 1$ iff the speaker is male;
 - and $\text{is_female} = \mathbb{1}_F = 1$ iff the speaker is female.
- The last option is encoded by setting **is_male** and **is_female** to 0.
- More generally, a nominal variable with K levels $(c_0, \dots, c_{K-1})$, can be “split” into $K - 1$ **binary variables**: $\mathbb{1}_{c_0}$ (=1 iff the class is c_0), $\mathbb{1}_{c_1}$ (=1 iff the class is c_1), ..., $\mathbb{1}_{c_{K-2}}$ (=1 iff the class is c_{K-2}).
- Why not include $\mathbb{1}_{c_{K-1}}$, the indicator function for the last class? This would be redundant: the last class is the class that is different from all the others; this can be encoded by setting $\mathbb{1}_{c_0}$, $\mathbb{1}_{c_1}$, ..., $\mathbb{1}_{c_{K-2}}$ to 0.

A full linear model of pause duration

$$\begin{aligned} \underbrace{\widehat{\text{pause}}}_{\text{continuous}} = & \beta_1 \underbrace{\text{length}}_{\text{discrete}} + \beta_2 \underbrace{\text{rate}}_{\text{continuous}} + \beta_3 \underbrace{\text{context}}_{\text{binary}} \\ & + \beta_4 \underbrace{\text{is_male}}_{\text{binary}} + \beta_5 \underbrace{\text{is_female}}_{\text{binary}} \\ & + \beta_6 \underbrace{\text{importance}}_{\text{discrete}} + \alpha \end{aligned}$$

- The `length` variable (measuring sentence length) is a positive integer. If its slope β_1 is negative, what does it mean about the influence of `length` on pause duration?
- The `is_female` variable is 1 iff the speaker is female. If its slope β_5 is positive, what does it mean about the influence of being female on pause duration?

A full linear model of pause duration

$$\begin{array}{ccccccc} \widehat{\text{pause}} & = & \beta_1 \text{length} & + & \beta_2 \text{rate} & + & \beta_3 \text{context} \\ \underbrace{\hspace{1.5cm}}_{\text{continuous}} & & \underbrace{\hspace{1.5cm}}_{\text{discrete}} & & \underbrace{\hspace{1.5cm}}_{\text{continuous}} & & \underbrace{\hspace{1.5cm}}_{\text{binary}} \\ & & & & + \beta_4 \text{is_male} & + & \beta_5 \text{is_female} \\ & & & & \underbrace{\hspace{1.5cm}}_{\text{binary}} & & \underbrace{\hspace{1.5cm}}_{\text{binary}} \\ & & & & & & + \beta_6 \text{importance} + \alpha \\ & & & & & & \underbrace{\hspace{1.5cm}}_{\text{discrete}} \end{array}$$

- The **length** variable (measuring sentence length) is a positive integer. If its slope β_1 is negative, what does it mean about the influence of **length** on pause duration?
- The **is_female** variable is 1 iff the speaker is female. If its slope β_5 is positive, what does it mean about the influence of being female on pause duration?

A full linear model of pause duration

$$\begin{array}{ccccccc} \widehat{\text{pause}} & = & \beta_1 \text{length} & + \beta_2 \text{rate} & + \beta_3 \text{context} & & \\ \underbrace{\hspace{1.5cm}}_{\text{continuous}} & & \underbrace{\hspace{1.5cm}}_{\text{discrete}} & \underbrace{\hspace{1.5cm}}_{\text{continuous}} & \underbrace{\hspace{1.5cm}}_{\text{binary}} & & \\ & & + \beta_4 \text{is_male} & + \beta_5 \text{is_female} & & & \\ & & \underbrace{\hspace{1.5cm}}_{\text{binary}} & \underbrace{\hspace{1.5cm}}_{\text{binary}} & & & \\ & & & + \beta_6 \text{importance} & + \alpha & & \\ & & & \underbrace{\hspace{1.5cm}}_{\text{discrete}} & & & \end{array}$$

- The **length** variable (measuring sentence length) is a positive integer. If its slope β_1 is negative, what does it mean about the influence of **length** on pause duration?
- The **is_female** variable is 1 iff the speaker is female. If its slope β_5 is positive, what does it mean about the influence of being female on pause duration?

Imbalances and the need for standardization

$$\underbrace{\widehat{\text{pause}}}_{\text{continuous}} = \beta_1 \underbrace{\text{length}}_{\text{discrete}} + \beta_2 \underbrace{\text{rate}}_{\text{continuous}} + \beta_3 \underbrace{\text{context}}_{\text{binary}} \\ + \beta_4 \underbrace{\text{is_male}}_{\text{binary}} + \beta_5 \underbrace{\text{is_female}}_{\text{binary}} \\ + \beta_6 \underbrace{\text{importance}}_{\text{discrete}} + \alpha$$

- Speech rate is typically encoded in words per minute (average 120-150 wpm).
- What if we chose to encode the **rate** as words per hour? per day? per millisecond?
- What would happen to the coefficient β_2 in each case?

Imbalances and the need for standardization

$$\underbrace{\widehat{\text{pause}}}_{\text{continuous}} = \beta_1 \underbrace{\text{length}}_{\text{discrete}} + \beta_2 \underbrace{\text{rate}}_{\text{continuous}} + \beta_3 \underbrace{\text{context}}_{\text{binary}} \\ + \beta_4 \underbrace{\text{is_male}}_{\text{binary}} + \beta_5 \underbrace{\text{is_female}}_{\text{binary}} \\ + \beta_6 \underbrace{\text{importance}}_{\text{discrete}} + \alpha$$

- Speech rate is typically encoded in words per minute (average 120-150 wpm).
- What if we chose to encode the **rate** as words per hour? per day? per millisecond?
- What would happen to the coefficient β_2 in each case?

Standardization

- Features are standardized to ensure they all vary on the same scale, i.e. they all have a mean of 0 and a standard deviation of 1. For each feature $k \in \{1, \dots, n\}$ and input i :

$$x_k^{(i)} := \frac{x_k^{(i)} - \text{MEAN}(x_k)}{\text{STD}(x_k)}$$

- Standardization is a linear operation, so it does not fundamentally affect the prediction – just the magnitude of the parameters (β s and α).
- Standardizing all independent variables makes their β s **directly comparable** to each other: they will be on the same scale. So the **relative influence of the different variables** can be assessed.
- This also makes the computation of an optimal set of parameters easier.

Standardization

- Features are standardized to ensure they all vary on the same scale, i.e. they all have a mean of 0 and a standard deviation of 1. For each feature $k \in \{1, \dots, n\}$ and input i :

$$x_k^{(i)} := \frac{x_k^{(i)} - \text{MEAN}(x_k)}{\text{STD}(x_k)}$$

- Standardization is a linear operation, so it does not fundamentally affect the prediction – just the magnitude of the parameters (β s and α).
- Standardizing all independent variables makes their β s **directly comparable** to each other: they will be on the same scale. So the **relative influence of the different variables** can be assessed.
- This also makes the computation of an optimal set of parameters easier.

Standardization

- Features are standardized to ensure they all vary on the same scale, i.e. they all have a mean of 0 and a standard deviation of 1. For each feature $k \in \{1, \dots, n\}$ and input i :

$$x_k^{(i)} := \frac{x_k^{(i)} - \text{MEAN}(x_k)}{\text{STD}(x_k)}$$

- Standardization is a linear operation, so it does not fundamentally affect the prediction – just the magnitude of the parameters (β s and α).
- Standardizing all independent variables makes their β s **directly comparable** to each other: they will be on the same scale. So the **relative influence of the different variables** can be assessed.
- This also makes the computation of an optimal set of parameters easier.

Standardization

- Features are standardized to ensure they all vary on the same scale, i.e. they all have a mean of 0 and a standard deviation of 1. For each feature $k \in \{1, \dots, n\}$ and input i :

$$x_k^{(i)} := \frac{x_k^{(i)} - \text{MEAN}(x_k)}{\text{STD}(x_k)}$$

- Standardization is a linear operation, so it does not fundamentally affect the prediction – just the magnitude of the parameters (β s and α).
- Standardizing all independent variables makes their β s **directly comparable** to each other: they will be on the same scale. So the **relative influence of the different variables** can be assessed.
- This also makes the computation of an optimal set of parameters easier.

Using and understanding a fitted regression model

Predictive use

- Linear regression is a **supervised** learning method: we learn the β s and α from known input-output pairs of the form (independent variables, dependent variable).
- Let's assume we did this and so our β s and α reflect the most accurate linear relation between the independent variables and the dependent variable.

$$\widehat{\text{pause}} = \beta_1^* \text{length} + \beta_2^* \text{rate} + \beta_3^* \text{context} \\ + \beta_4^* \text{is_male} + \beta_5^* \text{is_female} \\ + \beta_6^* \text{importance} + \alpha^*$$

- Once the model parameters (the β s and α) are properly optimized, we can use them to predict the dependent variable (**pause**) for new speech data, characterized by the independent variables **length**, **rate**, **context**, **gender**, and **importance**.

Predictive use

- Linear regression is a **supervised** learning method: we learn the β s and α from known input-output pairs of the form (independent variables, dependent variable).
- Let's assume we did this and so our β s and α reflect the most accurate linear relation between the independent variables and the dependent variable.

$$\widehat{\text{pause}} = \beta_1^* \text{length} + \beta_2^* \text{rate} + \beta_3^* \text{context} \\ + \beta_4^* \text{is_male} + \beta_5^* \text{is_female} \\ + \beta_6^* \text{importance} + \alpha^*$$

- Once the model parameters (the β s and α) are properly optimized, we can use them to predict the dependent variable (pause) for new speech data, characterized by the independent variables length, rate, context, gender, and importance.

Predictive use

- Linear regression is a **supervised** learning method: we learn the β s and α from known input-output pairs of the form (independent variables, dependent variable).
- Let's assume we did this and so our β s and α reflect the most accurate linear relation between the independent variables and the dependent variable.

$$\widehat{\text{pause}} = \beta_1^* \text{length} + \beta_2^* \text{rate} + \beta_3^* \text{context} \\ + \beta_4^* \text{is_male} + \beta_5^* \text{is_female} \\ + \beta_6^* \text{importance} + \alpha^*$$

- Once the model parameters (the β s and α) are properly optimized, we can use them to predict the dependent variable (**pause**) for new speech data, characterized by the independent variables **length**, **rate**, **context**, **gender**, and **importance**.

Hypothesis testing

- The **sign and magnitude** of the β s tell us how the independent variables influence the dependent variable!
- Additionally, one can test if any given β is statistically different from 0.¹ This tells us if the corresponding independent variable is a **significant predictor** of the dependent variable–i.e. if it's truly relevant.
- This is useful in **experimental settings**, e.g. when addressing the sociolinguistics question whether gender influences the duration of post-sentential pauses!
- Elaborations of the linear regression model, called **linear mixed effect models**, have become a standard way to test hypotheses in many disciplines, in CogSci and beyond.

¹The outcome will be dependent on the magnitude of β (high-magnitude β -estimates are more likely to be different from 0), but also on the size of the training set (bigger training sets yield more reliable estimates).

Hypothesis testing

- The **sign and magnitude** of the β s tell us how the independent variables influence the dependent variable!
- Additionally, one can test if any given β is statistically different from 0.¹ This tells us if the corresponding independent variable is a **significant predictor** of the dependent variable—i.e. if it's truly relevant.
- This is useful in **experimental settings**, e.g. when addressing the sociolinguistics question whether gender influences the duration of post-sentential pauses!
- Elaborations of the linear regression model, called **linear mixed effect models**, have become a standard way to test hypotheses in many disciplines, in CogSci and beyond.

¹The outcome will be dependent on the magnitude of β (high-magnitude β -estimates are more likely to be different from 0), but also on the size of the training set (bigger training sets yield more reliable estimates).

Hypothesis testing

- The **sign and magnitude** of the β s tell us how the independent variables influence the dependent variable!
- Additionally, one can test if any given β is statistically different from 0.¹ This tells us if the corresponding independent variable is a **significant predictor** of the dependent variable–i.e. if it's truly relevant.
- This is useful in **experimental settings**, e.g. when addressing the sociolinguistics question whether gender influences the duration of post-sentential pauses!
- Elaborations of the linear regression model, called **linear mixed effect models**, have become a standard way to test hypotheses in many disciplines, in CogSci and beyond.

¹The outcome will be dependent on the magnitude of β (high-magnitude β -estimates are more likely to be different from 0), but also on the size of the training set (bigger training sets yield more reliable estimates).


Hypothesis testing

- The **sign and magnitude** of the β s tell us how the independent variables influence the dependent variable!
- Additionally, one can test if any given β is statistically different from 0.¹ This tells us if the corresponding independent variable is a **significant predictor** of the dependent variable—i.e. if it's truly relevant.
- This is useful in **experimental settings**, e.g. when addressing the sociolinguistics question whether gender influences the duration of post-sentential pauses!
- Elaborations of the linear regression model, called **linear mixed effect models**, have become a standard way to test hypotheses in many disciplines, in CogSci and beyond.

¹The outcome will be dependent on the magnitude of β (high-magnitude β -estimates are more likely to be different from 0), but also on the size of the training set (bigger training sets yield more reliable estimates).

Link with neural nets

- Linear regression is **at the core of basic neural networks**.
- Roughly, basic “feed-forward” neural nets are stacked parallelized linear regressions, with some non-linearities added between them.²




Linear regression is a one-layer, one-unit neural net, without any activation function.

²We'll see one such non-linearity next week, in the form of the sigmoid function.

Link with neural nets

- Linear regression is **at the core of basic neural networks**.
- Roughly, basic “feed-forward” neural nets are stacked parallelized linear regressions, with some non-linearities added between them.²



Linear regression is a one-layer, one-unit neural net, without any activation function.

²We'll see one such non-linearity next week, in the form of the sigmoid function.

Bonus: fitting a linear model

Error for a single input-output pair

- Linear regression is a **supervised** learning method: we learn the β s and α from known input-output (IO) pairs of the form $(\underbrace{(x_1, \dots, x_n)}_{\text{independent variables}}, \underbrace{y}_{\text{dependent variable}})$. Let's focus on one such pair.
- Intuitively, we want our β s and α to minimize the **difference** between our estimate $\hat{y}$ and the observed value of y .
- What is used in practice (for statistical/analytical reasons), is the **square of the difference** between $\hat{y}$ and y .

$$(\hat{y} - y)^2 = (y - \hat{y})^2 = (\beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha - y)^2$$

- "Good" β s and α should lead to a $\hat{y}$ close to the real value y , and $(\hat{y} - y)^2$ should be small.

Error for a single input-output pair

- Linear regression is a **supervised** learning method: we learn the β s and α from known input-output (IO) pairs of the form $(\underbrace{(x_1, \dots, x_n)}_{\text{independent variables}}, \underbrace{y}_{\text{dependent variable}})$. Let's focus on one such pair.
- Intuitively, we want our β s and α to minimize the **difference** between our estimate $\hat{y}$ and the observed value of y .
- What is used in practice (for statistical/analytical reasons), is the **square of the difference** between $\hat{y}$ and y .

$$(\hat{y} - y)^2 = (y - \hat{y})^2 = (\beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha - y)^2$$

- "Good" β s and α should lead to a $\hat{y}$ close to the real value y , and $(\hat{y} - y)^2$ should be small.

Error for a single input-output pair

- Linear regression is a **supervised** learning method: we learn the β s and α from known input-output (IO) pairs of the form $(\underbrace{(x_1, \dots, x_n)}_{\text{independent variables}}, \underbrace{y}_{\text{dependent variable}})$. Let's focus on one such pair.
- Intuitively, we want our β s and α to minimize the **difference** between our estimate $\hat{y}$ and the observed value of y .
- What is used in practice (for statistical/analytical reasons), is the **square of the difference** between $\hat{y}$ and y .

$$(\hat{y} - y)^2 = (y - \hat{y})^2 = (\beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha - y)^2$$

- “Good” β s and α should lead to a $\hat{y}$ close to the real value y , and $(\hat{y} - y)^2$ should be small.

Error for a single input-output pair

- Linear regression is a **supervised** learning method: we learn the β s and α from known input-output (IO) pairs of the form $(\underbrace{(x_1, \dots, x_n)}_{\text{independent variables}}, \underbrace{y}_{\text{dependent variable}})$. Let's focus on one such pair.
- Intuitively, we want our β s and α to minimize the **difference** between our estimate $\hat{y}$ and the observed value of y .
- What is used in practice (for statistical/analytical reasons), is the **square of the difference** between $\hat{y}$ and y .

$$(\hat{y} - y)^2 = (y - \hat{y})^2 = (\beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha - y)^2$$

- “Good” β s and α should lead to a $\hat{y}$ close to the real value y , and $(\hat{y} - y)^2$ should be small.

The ordinary least square cost function

- In practice, we will fit our linear regression on **many IO pairs**. The i -th IO pair will bear superscript i : $\left(x_1^{(i)}, \dots, x_n^{(i)}, y^{(i)}\right)$. For a single IO pair, the error is then:

$$\left(\hat{y}^{(i)} - y^{(i)}\right)^2 = \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)}\right)^2$$

- We can average this error over all k IO pairs!

$$J(\beta_1, \dots, \beta_n, \alpha) = \frac{1}{k} \sum_{i=1}^k \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)}\right)^2$$

- $J(\beta_1, \dots, \beta_n, \alpha)$ is the **ordinary least-square (OLS) cost function**: it will be low if for all i , our estimate $\hat{y}^{(i)}$ is close to the real $y^{(i)}$.

The ordinary least square cost function

- In practice, we will fit our linear regression on **many IO pairs**. The i -th IO pair will bear superscript i : $\left(x_1^{(i)}, \dots, x_n^{(i)}\right), y^{(i)}$. For a single IO pair, the error is then:

$$\left(\hat{y}^{(i)} - y^{(i)}\right)^2 = \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)}\right)^2$$

- We can average this error over all k IO pairs!

$$J(\beta_1, \dots, \beta_n, \alpha) = \frac{1}{k} \sum_{i=1}^k \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)}\right)^2$$

- $J(\beta_1, \dots, \beta_n, \alpha)$ is the **ordinary least-square (OLS) cost function**: it will be low if for all i , our estimate $\hat{y}^{(i)}$ is close to the real $y^{(i)}$.

The ordinary least square cost function

- In practice, we will fit our linear regression on **many IO pairs**. The i -th IO pair will bear superscript i : $\left(x_1^{(i)}, \dots, x_n^{(i)}\right), y^{(i)}$. For a single IO pair, the error is then:

$$\left(\hat{y}^{(i)} - y^{(i)}\right)^2 = \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)}\right)^2$$

- We can average this error over all k IO pairs!

$$J(\beta_1, \dots, \beta_n, \alpha) = \frac{1}{k} \sum_{i=1}^k \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)}\right)^2$$

- $J(\beta_1, \dots, \beta_n, \alpha)$ is the **ordinary least-square (OLS) cost function**: it will be low if for all i , our estimate $\hat{y}^{(i)}$ is close to the real $y^{(i)}$.

Optimizing the OLS cost function

Goal: find $\beta_1, \dots, \beta_n$ and α that minimize $J(\beta_1, \dots, \beta_n, \alpha)$

- With jargon:

$$(\beta_1^*, \dots, \beta_n^*, \alpha^*) = \operatorname{argmin}_{(\beta_1, \dots, \beta_n, \alpha)} J(\beta_1, \dots, \beta_n, \alpha)$$

- Luckily, using **squared differences** between $\hat{y}^{(i)}$ and $y^{(i)}$ for all IO pair i provides us with an **analytical** solution to the above optimization problem! Analytical means the solution is given by a formula which is often straightforwardly computable.
- The proof and computation of this exact solution are quite involved though, but luckily we have libraries that do it for us!

Optimizing the OLS cost function

Goal: find $\beta_1, \dots, \beta_n$ and α that minimize $J(\beta_1, \dots, \beta_n, \alpha)$

- With jargon:

$$(\beta_1^*, \dots, \beta_n^*, \alpha^*) = \operatorname{argmin}_{(\beta_1, \dots, \beta_n, \alpha)} J(\beta_1, \dots, \beta_n, \alpha)$$

- Luckily, using **squared differences** between $\hat{y}^{(i)}$ and $y^{(i)}$ for all IO pair i provides us with an **analytical** solution to the above optimization problem! Analytical means the solution is given by a formula which is often straightforwardly computable.
- The proof and computation of this exact solution are quite involved though, but luckily we have libraries that do it for us!

Optimizing the OLS cost function

Goal: find $\beta_1, \dots, \beta_n$ and α that minimize $J(\beta_1, \dots, \beta_n, \alpha)$

- With jargon:

$$(\beta_1^*, \dots, \beta_n^*, \alpha^*) = \operatorname{argmin}_{(\beta_1, \dots, \beta_n, \alpha)} J(\beta_1, \dots, \beta_n, \alpha)$$

- Luckily, using **squared differences** between $\hat{y}^{(i)}$ and $y^{(i)}$ for all IO pair i provides us with an **analytical** solution to the above optimization problem! Analytical means the solution is given by a formula which is often straightforwardly computable.
- The proof and computation of this exact solution are quite involved though, but luckily we have libraries that do it for us!

Optimizing the OLS cost function

Goal: find $\beta_1, \dots, \beta_n$ and α that minimize $J(\beta_1, \dots, \beta_n, \alpha)$

- With jargon:

$$(\beta_1^*, \dots, \beta_n^*, \alpha^*) = \operatorname{argmin}_{(\beta_1, \dots, \beta_n, \alpha)} J(\beta_1, \dots, \beta_n, \alpha)$$

- Luckily, using **squared differences** between $\hat{y}^{(i)}$ and $y^{(i)}$ for all IO pair i provides us with an **analytical** solution to the above optimization problem! Analytical means the solution is given by a formula which is often straightforwardly computable.
- The proof and computation of this exact solution are quite involved though, but luckily we have libraries that do it for us!

Assessing goodness-of-fit

- A basic linear regression model will fit the β s and α , and return a R^2 (**R-squared**) score.
- The R^2 is between 0 (bad) and 1 (good) and measures the proportion of the variation in the dependent variable that is predictable from the independent variable(s) – in brief, **how much variance our model explains; how well we fit**.
- What's a “good” R^2 depends on the discipline.
 - The threshold will be higher (.7-.9) in hard sciences (for which the factors are better known, and precisely measurable);
 - and lower (.3-.5) in social sciences (for which the factors are multifaceted, and sometimes noisy).

Assessing goodness-of-fit

- A basic linear regression model will fit the β s and α , and return a R^2 (**R-squared**) score.
- The R^2 is between 0 (bad) and 1 (good) and measures the proportion of the variation in the dependent variable that is predictable from the independent variable(s) – in brief, **how much variance our model explains; how well we fit**.
- What's a “good” R^2 depends on the discipline.
 - The threshold will be higher (.7-.9) in hard sciences (for which the factors are better known, and precisely measurable);
 - and lower (.3-.5) in social sciences (for which the factors are multifaceted, and sometimes noisy).

Assessing goodness-of-fit

- A basic linear regression model will fit the β s and α , and return a R^2 (**R-squared**) score.
- The R^2 is between 0 (bad) and 1 (good) and measures the proportion of the variation in the dependent variable that is predictable from the independent variable(s) – in brief, **how much variance our model explains; how well we fit**.
- What's a “good” R^2 depends on the discipline.
 - The threshold will be higher (.7-.9) in hard sciences (for which the factors are better known, and precisely measurable);
 - and lower (.3-.5) in social sciences (for which the factors are multifaceted, and sometimes noisy).

Assessing goodness-of-fit

- A basic linear regression model will fit the β s and α , and return a R^2 (**R-squared**) score.
- The R^2 is between 0 (bad) and 1 (good) and measures the proportion of the variation in the dependent variable that is predictable from the independent variable(s) – in brief, **how much variance our model explains; how well we fit**.
- What's a “good” R^2 depends on the discipline.
 - The threshold will be higher (.7-.9) in hard sciences (for which the factors are better known, and precisely measurable);
 - and lower (.3-.5) in social sciences (for which the factors are multifaceted, and sometimes noisy).

Assessing goodness-of-fit

- A basic linear regression model will fit the β s and α , and return a R^2 (**R-squared**) score.
- The R^2 is between 0 (bad) and 1 (good) and measures the proportion of the variation in the dependent variable that is predictable from the independent variable(s) – in brief, **how much variance our model explains; how well we fit**.
- What's a “good” R^2 depends on the discipline.
 - The threshold will be higher (.7-.9) in hard sciences (for which the factors are better known, and precisely measurable);
 - and lower (.3-.5) in social sciences (for which the factors are multifaceted, and sometimes noisy).

Fixes and improvements

Crazy coefficients

- Suppose we predicting **pause** only based on **context** (formal/informal) and **importance** (unimportant/important):

$$\widehat{\text{pause}} = \beta_1 \text{context} + \beta_2 \text{importance} + \alpha$$

- Suppose that in our training data **context** and **importance** are perfectly correlated: formal contexts always involve important topics, and informal contexts, unimportant topics. We'll train our model on data that look like:

$$y^{(1)} = \beta_1 \times 0 + \beta_2 \times 1 + \alpha = \beta_1 + \alpha \quad \text{formal/important features}$$

$$y^{(2)} = \beta_1 \times 1 + \beta_2 \times 0 + \alpha = \beta_2 + \alpha \quad \text{informal/unimportant features}$$

- Issue: if $(\beta_1^*, \beta_2^*, \alpha^*)$ is optimal, so is $(\beta_1^* - 1000, \beta_2^* - 1000, \alpha^* + 1000)$, but these 2 sets of parameters make very different predictions for possible informal/important and formal/unimportant features bundles...

Crazy coefficients

- Suppose we predicting **pause** only based on **context** (formal/informal) and **importance** (unimportant/important):

$$\widehat{\text{pause}} = \beta_1 \text{context} + \beta_2 \text{importance} + \alpha$$

- Suppose that in our training data **context** and **importance** are perfectly correlated: formal contexts always involve important topics, and informal contexts, unimportant topics. We'll train our model on data that look like:

$y^{(1)}$	$=$	$\beta_1 \times 0 + \beta_2 \times 1 + \alpha = \beta_1 + \alpha$	formal/important features
$y^{(2)}$	$=$	$\beta_1 \times 1 + \beta_2 \times 0 + \alpha = \beta_2 + \alpha$	informal/unimportant features

- Issue: if $(\beta_1^*, \beta_2^*, \alpha^*)$ is optimal, so is $(\beta_1^* - 1000, \beta_2^* - 1000, \alpha^* + 1000)$, but these 2 sets of parameters make very different predictions for possible informal/important and formal/unimportant features bundles...

Crazy coefficients

- Suppose we predicting **pause** only based on **context** (formal/informal) and **importance** (unimportant/important):

$$\widehat{\text{pause}} = \beta_1 \text{context} + \beta_2 \text{importance} + \alpha$$

- Suppose that in our training data **context** and **importance** are perfectly correlated: formal contexts always involve important topics, and informal contexts, unimportant topics. We'll train our model on data that look like:

$$\begin{aligned} y^{(1)} &= \beta_1 \times 0 + \beta_2 \times 1 + \alpha = \beta_1 + \alpha && \text{formal/important features} \\ y^{(2)} &= \beta_1 \times 1 + \beta_2 \times 0 + \alpha = \beta_2 + \alpha && \text{informal/unimportant features} \end{aligned}$$

- Issue: if $(\beta_1^*, \beta_2^*, \alpha^*)$ is optimal, so is $(\beta_1^* - 1000, \beta_2^* - 1000, \alpha^* + 1000)$, but these 2 sets of parameters make very different predictions for possible informal/important and formal/unimportant features bundles...

Regularization

- To prevent the β s from turning crazy when the training data is a little skewed, we add a regularization term to the regression formula, **penalizing high-magnitude β s**.
- This term includes a **hyperparameter** λ , which quantifies how much regularization should influence the prediction.³

$$\hat{y} = \underbrace{\beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha}_{\text{basic linear regression model}} + \underbrace{\lambda \sum_{j=1}^n \beta_j^2}_{\text{regularization term}}$$

$$\widehat{\text{pause}} = \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} \\ + \beta_4 \text{is_male} + \beta_5 \text{is_female} + \beta_6 \text{importance} + \alpha \\ + \lambda(\beta_1^2 + \beta_2^2 + \beta_3^2 + \beta_4^2 + \beta_5^2 + \beta_6^2)$$

³Different values of λ can be tried on different test sets, and the one minimizing the loss J can be eventually selected.

Regularization

- To prevent the β s from turning crazy when the training data is a little skewed, we add a regularization term to the regression formula, **penalizing high-magnitude β s**.
- This term includes a **hyperparameter** λ , which quantifies how much regularization should influence the prediction.³

$$\hat{y} = \underbrace{\beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \alpha}_{\text{basic linear regression model}} + \underbrace{\lambda \sum_{j=1}^n \beta_j^2}_{\text{regularization term}}$$

$$\begin{aligned} \widehat{\text{pause}} = & \beta_1 \text{length} + \beta_2 \text{rate} + \beta_3 \text{context} \\ & + \beta_4 \text{is_male} + \beta_5 \text{is_female} + \beta_6 \text{importance} + \alpha \\ & + \lambda(\beta_1^2 + \beta_2^2 + \beta_3^2 + \beta_4^2 + \beta_5^2 + \beta_6^2) \end{aligned}$$

³Different values of λ can be tried on different test sets, and the one minimizing the loss J can be eventually selected.

Advantages of regularization

- Adding a regularization term to the linear model affects the corresponding OLS function! It now includes the exact same regularization term.

$$J(\beta_1, \dots, \beta_n, \alpha) = \lambda \sum_{j=1}^n \beta_j^2 + \frac{1}{k} \sum_{i=1}^k \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)} \right)^2$$

- To properly minimize this new OLS function, one must keep $\lambda \sum_{j=1}^n \beta_j^2$ low, meaning, one must keep the β s low in magnitude!
- Regularization prevents **overfitting**, in particular a model making extreme predictions for unseen inputs.
- Interestingly, it also makes the computation of an optimal solution $(\beta_1^*, \dots, \beta_n^*, \alpha^*)$ easier!

Advantages of regularization

- Adding a regularization term to the linear model affects the corresponding OLS function! It now includes the exact same regularization term.

$$J(\beta_1, \dots, \beta_n, \alpha) = \lambda \sum_{j=1}^n \beta_j^2 + \frac{1}{k} \sum_{i=1}^k \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)} \right)^2$$

- To properly minimize this new OLS function, one must keep $\lambda \sum_{j=1}^n \beta_j^2$ low, meaning, one must keep the β s low in magnitude!
- Regularization prevents **overfitting**, in particular a model making extreme predictions for unseen inputs.
- Interestingly, it also makes the computation of an optimal solution $(\beta_1^*, \dots, \beta_n^*, \alpha^*)$ easier!

Advantages of regularization

- Adding a regularization term to the linear model affects the corresponding OLS function! It now includes the exact same regularization term.

$$J(\beta_1, \dots, \beta_n, \alpha) = \lambda \sum_{j=1}^n \beta_j^2 + \frac{1}{k} \sum_{i=1}^k \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)} \right)^2$$

- To properly minimize this new OLS function, one must keep $\lambda \sum_{j=1}^n \beta_j^2$ low, meaning, one must keep the β s low in magnitude!
- Regularization prevents **overfitting**, in particular a model making extreme predictions for unseen inputs.
- Interestingly, it also makes the computation of an optimal solution $(\beta_1^*, \dots, \beta_n^*, \alpha^*)$ easier!

Advantages of regularization

- Adding a regularization term to the linear model affects the corresponding OLS function! It now includes the exact same regularization term.

$$J(\beta_1, \dots, \beta_n, \alpha) = \lambda \sum_{j=1}^n \beta_j^2 + \frac{1}{k} \sum_{i=1}^k \left(\beta_1 x_1^{(i)} + \beta_2 x_2^{(i)} + \dots + \beta_n x_n^{(i)} + \alpha - y^{(i)} \right)^2$$

- To properly minimize this new OLS function, one must keep $\lambda \sum_{j=1}^n \beta_j^2$ low, meaning, one must keep the β s low in magnitude!
- Regularization prevents **overfitting**, in particular a model making extreme predictions for unseen inputs.
- Interestingly, it also makes the computation of an optimal solution $(\beta_1^*, \dots, \beta_n^*, \alpha^*)$ easier!

Appendix

Thermometer encoding

- Suppose we have an ordinal variable whose levels are **not evenly distributed**. Perhaps we are interested in how emotionally affected the speaker is, and want to be fine-grained about extremes: affected/very affected/extremely affected.
- In that case, mapping the levels to integers would wrongly suggest differences between levels are the same!
- Instead, we use **a series of indicator functions, that we see as a “thermometer”**. Our functions are $\mathbb{1}_{\leq \text{affected}}$, $\mathbb{1}_{\leq \text{v.affected}}$, $\mathbb{1}_{\leq \text{v.v.affected}}$.
 - To encode an affected mood, we set $\mathbb{1}_{\leq \text{affected}}$ to 1 and the 2 other variables to 0 ($\simeq$ low temperature thermometer).
 - To encode a very affected mood, we set $\mathbb{1}_{\leq \text{affected}}$ and $\mathbb{1}_{\leq \text{v.affected}}$ to 1 and $\mathbb{1}_{\leq \text{v.v.affected}}$ to 0 ($\simeq$ mid temperature thermometer).
 - To encode an extremely affected mood, we set all variables to 1 ($\simeq$ high temperature thermometer)
- The number of variables set to 1 is like the level of a thermometer. But the differences between levels are qualitative instead of quantitative.