

LIN7076 – Foundations of Computational Linguistics


N-gram models

Adèle Hénot-Mortier

28/09/2026


Queen Mary University of London

Generating text



The intuition behind n-grams

- Your next word depends on the words you have used before.
- When predicting a word, pick the most likely, **given what was written before, within a specific window.**



- Likelihood (the ps) gets approximated by computing **statistics over a similar-enough training corpus.**

Plan for today

- Understand the n -gram **prediction process**.
- Understand how the probabilities used in that process are **approximated** from a training corpus.
- Understand how the resulting model can be **evaluated** against a test corpus.
- Explore **fixes/improvements** of shortcomings of the basic model.

The n-gram generation process

Why n-gram?

When predicting a word, pick the **most likely, given what was written before** within a specific window.

- We want to compute the conditional probability of a word w_i , given a bunch of preceding words $[w_{i-n+1}, \dots, w_{i-1}]$, that we call the **context**, of size $n - 1$.
- The sequence $[w_{i-n+1}, \dots, w_{i-1}] + [w_i]$ is of size n . We call it a **n -gram**.
 - For instance, the sequence $[I, had, such, a]$, is a 4-gram.
- A **n -gram model**, predicts the next word, given a $(n - 1)$ -gram of preceding words.
 - For instance, a 4-gram model can predict a to follow the 3-gram $[I, had, such]$.


Counting ns

- What's the size of the sequence [The blue book]? 3! What kind of n -gram is it? A 3-gram!
- Suppose you want to predict the word following [The blue book]. What kind of n -gram model should you use? A 4-gram model, because I want to predict the next word based on a context window of size 3 ($= 4 - 1$). Put differently, it's because my sequence *after* I predict the next word will be of size 4.
- Suppose you want to predict the word following [book] (you “forgot” that *book* is preceded by *the* and *blue*). What kind of n -gram model should you use? A 2-gram (“bigram”) model, because I want to predict the next word based on a context window of size 1 ($= 2 - 1$). Put differently, it's because my sequence *after* I predict the next word will be of size 2.

Refresher on conditional probabilities

- The **conditional probability** of an event A , given an event B , is the probability of A when B is known to happen.
- It can be computed using the joint probability of A and B , and B 's probability.

$$\begin{array}{ccccc} \text{conditional} & & \text{joint probability} & & \text{probability} \\ \text{probability} & & \text{of } A \text{ and } B & & \text{A and B} \\ \text{of } B \text{ given } A & & & & \text{occur together} \\ \underbrace{P(A|B)} & = & \frac{\overbrace{P(A, B)}}{P(B)} & = & \frac{\overbrace{P(A \cap B)}}{P(B)} \end{array}$$



- Given any fixed event B , the probabilities $\{P(X|B) \mid X\}$, form a probability distribution (i.e. are each in $[0, 1]$ and sum to 1)!

Quick example

- Suppose our sentence is a polar (=yes-no) question and our first word is either *Are* (with probability 25%) or *Is* (with probability 75%).
- Suppose moreover that [*Are the*] is a sequence with 10% probability and [*Is the*] a sequence with 15% probability.
- What's the conditional probability of *the* given *Are*? Here, *A* is the event that the second word is *the*, while *B* is the event that the first word is *Are*, and we compute
$$P(A|B) = P(\text{Are the})/P(\text{Are}) = .1/.25 = .4$$
- What's the conditional probability of *the* given *Is*? Here, *A* is the event that the second word is *the*, while *B* is the event that the first word is *Is*, and we compute
$$P(A|B) = P(\text{Is the})/P(\text{Is}) = .15/.75 = .2$$

Rephrasing the n-gram intuition using conditional probabilities

- Our target event A , is that of the word w_i being used. We call this word **target**.
- Our given event B , is that of the words $[w_{i-n+1}, \dots, w_{i-1}]$ being used before. We call this sequence **context**.
- We want to compute $P(A|B)$, the probability of w_i given $[w_{i-n+1}, \dots, w_{i-1}]$!

Rephrasing the n-gram intuition using conditional probabilities

- $P(w_i | [w_{i-n+1}, \dots, w_{i-1}])$ depends solely on the context, and the context appended with the target (= the n-gram!).

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

$$\begin{aligned} P(w_i | [w_{i-n+1}, \dots, w_{i-1}]) &= \frac{P([w_{i-n+1}, \dots, w_i])}{P([w_{i-n+1}, \dots, w_{i-1}])} \\ &= \frac{P(\text{context} + \text{target})}{P(\text{context})} \end{aligned}$$

$$P(\text{great} | \text{I had such a}) = \frac{P(\text{I had such a great})}{P(\text{I had such a})}$$

- What would a **unigram model** look like? In a unigram model, we predict the next word based on a context of size 0 (=1 – 1). So our conditional probability is $P(w_i | []) = P(w_i)$ (sometimes called “prior”). It is fully oblivious of the left context.

A 3-gram example with made-up probabilities

- Suppose our sentence starts with the bigram *I had*. Suppose that *such* and *some* are the only 2 possible continuations.

$$P(\text{such} | \text{I had}) = \frac{P(\text{I had such})}{P(\text{I had})} = .7$$

$$P(\text{some} | \text{I had}) = \frac{P(\text{I had some})}{P(\text{I had})} = .3$$

- *Such* is more likely so we pick *such*. Now our sentence is *I had such*, and our context (bigram) becomes *had such*. Suppose that now *a* and *good* are the only 2 possible continuations.

$$P(a | \text{had such}) = \frac{P(\text{had such a})}{P(\text{had such})} = .8$$

$$P(\text{good} | \text{had such}) = \frac{P(\text{had such good})}{P(\text{had such})} = .2$$

- *A* is more likely, so we pick *a*. Now our sentence is *I had such a*, and our context becomes *such a*. Etc...

A potential issue with fixed window sizes

When predicting a word, pick the most likely, given what was written before **within a specific window**.

- Why use a window, especially a short one? Natural language is recursive after all!
- Consider (1), which features **center-embedding**, in the form of an arbitrary number of nested complementizer phrases, between the main subject and its verb..¹

(1) The *student* who Jo said that Ed thinks that ... Al likes __ ...
subject arbitrarily long, nested relative clauses

- Could a n -gram model with a fixed window size correctly predict the number marking on the following verb? Could it even predict that it should be a verb?

¹You can try to do the same with only relative clauses!

Why windows, anyway

- Using a fixed window makes the assumption that most of the time, language generation is oblivious of what was said too long ago. As we've just seen, this is not always true.
- This is known as the **Markov assumption**, which applies to processes well beyond CompLing! Biology, economy, robotics...
- Even if imperfect, the Markov assumption is computationally more efficient, as it requires us to compute n -gram and $n - 1$ -gram **occurrences** for a fixed n , as we will soon see.
- It will also prevent the model from “memorizing” big chunks of the training data, something known as overfitting.
- Can you think of what shorter/longer windows would on average best capture in terms of syntax/semantics?

A note on sampling

When predicting a word, **pick the most likely**, given what was written before within a specific window.

- What if two possible targets w_i and w'_i are associated to very close conditional probabilities, e.g. .999 and .1? Should we still pick the most likely, all the time?
- In practice, we **sample**: we pick target words randomly, giving more chances to the more likely ones.
- One simple way to think about it: drawing one ball from a big lottery box such that each target word is assigned a ball color, and gets as many balls as its conditional probability converted to a round percentage (so $p = .155$ means 16 balls).



Sampling example

- Remember our made up conditional probabilities for *such* and *some*, given the context *I had*?

$$P(\text{such} | I \text{ had}) = \frac{P(I \text{ had such})}{P(I \text{ had})} = .7$$
$$P(\text{some} | I \text{ had}) = \frac{P(I \text{ had some})}{P(I \text{ had})} = .3$$

- Instead of returning *such* (presumed most likely) 100% of the time, sampling will return *such* 70% of the time, and *some* 30% of the time.
- More generally, any target word *w* would be picked $(100 \times P(w | I \text{ had}))\%$ of the time.

Approximating conditional probabilities

From probabilities to counts

- So far we have reasoned about hypothetical probabilities of words or sequences of words. **How are these probabilities estimated?**
- The probability of a sequence of words S of any fixed length (can be 1, i.e. S is a single word) can be estimated based on a training corpus.
- This is done by counting the number of S 's occurrences, and dividing it up by the sum of the counts of all same-length sequences.
- The [Google \$n\$ -gram Viewer](#) actually does this on a very large corpus of books!

From probabilities to counts

- [The following demonstration is a bit advanced; you will not be asked to reproduce or use it in assessments]. Let $C(S)$ be the number of occurrences of S in a given corpus.

$$P(S) = \frac{C(S)}{\sum_{|S'|=|S|} C(S')} \triangleq \alpha C(S)$$

- Suppose now S' is obtained from S by removing S 's last word: $S = S' + w$. If S is a n -gram, S' is a $(n - 1)$ -gram. S' 's probability can be expressed by summing the probabilities of its immediate continuations...

$$P(S') = \sum_w P(S' + w) = \sum_w \alpha C(S' + w) = \alpha \sum_w C(S' + w) = \alpha C(S')$$

- We have just expressed probabilities of n and $(n - 1)$ -grams using their **counts** and the **same normalizing factor** α !

From conditional probabilities to counts

$$\begin{aligned}P([w_{i-n+1}, \dots, w_i]) &= \alpha C([w_{i-n+1}, \dots, w_i]) \\P([w_{i-n+1}, \dots, w_{i-1}]) &= \alpha C([w_{i-n+1}, \dots, w_{i-1}])\end{aligned}$$

- Plugging this into our n -gram prediction formula, the α s eventually cancel out! So we get a conditional probability that solely depends on n -gram counts.

$$\begin{aligned}P(w_i | [w_{i-n+1}, \dots, w_{i-1}]) &= \frac{P([w_{i-n+1}, \dots, w_i])}{P([w_{i-n+1}, \dots, w_{i-1}])} \\&\approx \frac{C([w_{i-n+1}, \dots, w_i])}{C([w_{i-n+1}, \dots, w_{i-1}])} \\&\approx \frac{C(\text{context} + \text{target})}{C(\text{context})}\end{aligned}$$

Quick example

- Suppose our corpus only contains polar questions whose first word is either *Are* (1050 occurrences) or *Is* (3150 occurrences).
- Suppose moreover that [*Are the*] is a sequence occurring 420 times sentence-initially in the corpus and [*Is the*] occurs 630 times, also sentence-initially.
- What's the conditional probability of *the* given *Are* approximated from that corpus? We compute
$$P(\textit{the}|\textit{Are}) \simeq c(\textit{Are the})/c(\textit{Are}) = 420/1050 = .4$$
- What's the conditional probability of *the* given *Is* approximated from that corpus? We compute
$$P(\textit{the}|\textit{Is}) \simeq c(\textit{Is the})/c(\textit{Is}) = 630/3150 = .2$$
- These conditional probabilities should look familiar: it's because I made the count s.t. they correspond to the probabilities of the previous exercise $\times 4200$ (i.e. $\alpha = 1/4200$).

Dealing with sentence boundaries

$$P(w_i | [w_{i-n+1}, \dots, w_{i-1}]) \simeq \frac{C([w_{i-n+1}, \dots, w_i])}{C([w_{i-n+1}, \dots, w_{i-1}])} = \frac{C(\text{context} + \text{target})}{C(\text{context})}$$

- The above equation predicts the next word, given a context of size $n - 1$. But what should we do at the very beginning of a sentence, when the context is empty? And when should we stop?
- To properly initiate and terminate generation, we posit two extra silent “words”: **<s>** (“**beginning** of sentence”) and **</s>** (“**end** of sentence”).
- For a n -gram model, each sentence in the training corpus gets **padded with $n - 1$ <s> at the beginning, and 1 </s> at the end.** Counts incorporate these extra words.
- To start generating a sentence, we can then assume an initial, silent context made of $n - 1$ <s> symbols!
- And we stop as soon as </s> gets generated.

Estimating conditional probabilities: a worked example

- Let's compute a toy bigram model. We pad with 1 `<s>` symbol.

(2) a. `<s>` my dog is nice `</s>`

b. `<s>` my cat is nasty `</s>`

c. `<s>` your cat hates my cat `</s>`

- Bigrams: (`<s>`, my), (my, dog), (dog, is), (is, nice), (nice, `</s>`), (`<s>`, my), (my, cat), (cat, is), (is, nasty), (nasty, `</s>`), (`<s>`, Your), (Your, cat), (cat, hates), (hates, my), (my, cat), (cat, `</s>`).
- Unigrams: `<s>`, my, dog, is, nice, `</s>`, `<s>`, my, cat, is, nasty, `</s>`, `<s>`, Your, cat, hates, my, cat, `</s>`.

Estimating conditional probabilities: a worked example

- (3) a. $\langle s \rangle$ my dog is nice $\langle /s \rangle$
b. $\langle s \rangle$ my cat is nasty $\langle /s \rangle$
c. $\langle s \rangle$ your cat hates my cat $\langle /s \rangle$

$w_i \rightarrow$ $\downarrow w_{i-1}$	my	dog	is	nice	cat	nasty	your	hates	$\langle /s \rangle$	Total
$\langle s \rangle$	2	0	0	0	0	0	1	0	0	3
my	0	1	0	0	2	0	0	0	0	3
dog	0	0	1	0	0	0	0	0	0	1
is	0	0	0	1	0	1	0	0	0	2
nice	0	0	0	0	0	0	0	0	1	1
cat	0	0	1	0	0	0	0	1	1	3
nasty	0	0	0	0	0	0	0	0	1	1
your	0	0	0	0	1	0	0	0	0	1
hates	1	0	0	0	0	0	0	0	0	1
Total	3	1	2	1	3	1	1	1	3	

$$P(\text{my}|\text{<s>}) = \frac{C(\text{<s>my})}{C(\text{<s>})} = \frac{2}{3}$$

$w_i \rightarrow$ $\downarrow w_{i-1}$	my	dog	is	nice	cat	nasty	your	hates	</s>	Total
<s>	2	0	0	0	0	0	1	0	0	3
my	0	1	0	0	2	0	0	0	0	3
dog	0	0	1	0	0	0	0	0	0	1
is	0	0	0	1	0	1	0	0	0	2
nice	0	0	0	0	0	0	0	0	1	1
cat	0	0	1	0	0	0	0	1	1	3
nasty	0	0	0	0	0	0	0	0	1	1
your	0	0	0	0	1	0	0	0	0	1
hates	1	0	0	0	0	0	0	0	0	1
Total	3	1	2	1	3	1	1	1	3	

$$P(\text{dog}|\text{my}) = \frac{C(\text{my dog})}{C(\text{my})} = \frac{1}{3}$$

$w_i \rightarrow$ $\downarrow w_{i-1}$	my	dog	is	nice	cat	nasty	your	hates	</s>	Total
<s>	2	0	0	0	0	0	1	0	0	3
my	0	1	0	0	2	0	0	0	0	3
dog	0	0	1	0	0	0	0	0	0	1
is	0	0	0	1	0	1	0	0	0	2
nice	0	0	0	0	0	0	0	0	1	1
cat	0	0	1	0	0	0	0	1	1	3
nasty	0	0	0	0	0	0	0	0	1	1
your	0	0	0	0	1	0	0	0	0	1
hates	1	0	0	0	0	0	0	0	0	1
Total	3	1	2	1	3	1	1	1	3	

$$P(\text{hates}|\text{cat}) = \frac{C(\text{cat hates})}{C(\text{cat})} = \frac{1}{3}$$

$w_i \rightarrow$ $\downarrow w_{i-1}$	my	dog	is	nice	cat	nasty	your	hates	</s>	Total
<s>	2	0	0	0	0	0	1	0	0	3
my	0	1	0	0	2	0	0	0	0	3
dog	0	0	1	0	0	0	0	0	0	1
is	0	0	0	1	0	1	0	0	0	2
nice	0	0	0	0	0	0	0	0	1	1
cat	0	0	1	0	0	0	0	1	1	3
nasty	0	0	0	0	0	0	0	0	1	1
your	0	0	0	0	1	0	0	0	0	1
hates	1	0	0	0	0	0	0	0	0	-
Total	3	1	2	1	3	1	1	1	3	

- Can you guess which sentence(s) would be most likely to be generated by this bigram model?

Bonus: Evaluating n-gram models

Probability of a corpus

- To evaluate how good a n -gram model is, one needs new data: typically a test corpus that is not too different from the training corpus.²
- We can then evaluate how likely this whole corpus is to be generated by our n -gram model!
- This is thanks to the chain rule:

$$\begin{aligned} P(A_1 \cap A_2 \cap \dots \cap A_k) &= P(A_1) \times \prod_{i=2}^k P(A_i | A_1 \cap \dots \cap A_{i-1}) \\ &\stackrel{\text{Markov}}{=} P(A_1) \times \prod_{i=2}^k P(A_i | A_{i-n+1} \cap \dots \cap A_{i-1}) \\ P([w_1 \dots w_k]) &\stackrel{\text{Markov}}{=} P(w_1) \times \prod_{i=2}^k P(w_i | [w_{n-i+1} \dots w_{i-1}]) \end{aligned}$$

²Why not test on the training corpus?

From probability to perplexity

- Measuring the performance of a n -gram based on the probability of a test corpus is nice, but depends on the corpus size: bigger corpora will mechanically be less probable (why?).
- To avoid this issue, we use a normalized variant of corpus probability called **perplexity**. The lower the perplexity, the better the model.

$$\begin{aligned} PPL([w_1 \dots w_k]) &= P([w_1 \dots w_k])^{-\frac{1}{k}} \\ &\stackrel{\text{Markov}}{=} \left(P(w_1) \times \prod_{i=2}^k P(w_i | [w_{n-i+1} \dots w_{i-1}]) \right)^{-\frac{1}{k}} \end{aligned}$$

- If $k \uparrow$, $\frac{1}{k} \downarrow$, $-\frac{1}{k} \uparrow$ and $p^{-\frac{1}{k}} \downarrow$ (with $p \geq 0$). So if two corpora have same probability, the larger one will get a smaller perplexity.

An issue with perplexity

$$PPL([w_1 \dots w_k]) \stackrel{\text{Markov}}{=} \left(P(w_1) \times \prod_{i=2}^k P(w_i | [w_{n-i+1} \dots w_{i-1}]) \right)^{-\frac{1}{k}}$$

- Looking at the above product of probabilities, we notice that if only one probability is 0, the whole product is $\frac{1}{0} = +\infty$!
- This would happen if the test corpus contains **one n -gram that the training corpus does not have** – which is far from unlikely!
- So perplexity is **unlikely to distinguish good from bad n -gram models**, as we defined them. Both kinds will likely get an infinite perplexity.
- More generally, probabilities learned from a finite training corpus will necessarily be 0 for many perfectly reasonable n -grams...bringing us to the **competence-performance distinction**.

A solution: smoothing

- To prevent perplexity from vanishing, we can simply **add a small increment** to all our n -gram counts!

$w_i \rightarrow$ $\downarrow w_{i-1}$	<s>	my	dog	is	nice	cat	nasty	your	hates	</s>	Total
<s>	ϵ	$2+\epsilon$	ϵ	ϵ	ϵ	ϵ	ϵ	$1+\epsilon$	ϵ	ϵ	$3+10\epsilon$
my	ϵ	ϵ	$1+\epsilon$	ϵ	ϵ	$2+\epsilon$	ϵ	ϵ	ϵ	ϵ	$3+10\epsilon$
dog	ϵ	ϵ	ϵ	$1+\epsilon$	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	$1+10\epsilon$
is	ϵ	ϵ	ϵ	ϵ	$1+\epsilon$	ϵ	$1+\epsilon$	ϵ	ϵ	ϵ	$2+10\epsilon$
nice	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	$1+\epsilon$	$1+10\epsilon$
cat	ϵ	ϵ	ϵ	$1+\epsilon$	ϵ	ϵ	ϵ	ϵ	$1+\epsilon$	$1+\epsilon$	$3+10\epsilon$
nasty	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	$1+\epsilon$	$1+10\epsilon$
your	ϵ	ϵ	ϵ	ϵ	ϵ	$1+\epsilon$	ϵ	ϵ	ϵ	ϵ	$1+10\epsilon$
hates	ϵ	$1+\epsilon$	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	$1+10\epsilon$
</s>	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	ϵ	-
Total	-	$3+10\epsilon$	$1+10\epsilon$	$2+10\epsilon$	$1+10\epsilon$	$3+10\epsilon$	$1+10\epsilon$	$1+10\epsilon$	$1+10\epsilon$	$3+10\epsilon$	

- How do you think increasing/decreasing ϵ will affect training set perplexity? Test set perplexity?

How smoothing affects conditional probabilities

$$\begin{aligned} P(w_i | [w_{i-n+1}, \dots, w_{i-1}]) &\simeq \frac{C_{\text{smooth}}([w_{i-n+1}, \dots, w_i])}{C_{\text{smooth}}([w_{i-n+1}, \dots, w_{i-1}])} \\ &\simeq \frac{C([w_{i-n+1}, \dots, w_i]) + \epsilon}{C([w_{i-n+1}, \dots, w_{i-1}]) + \epsilon|L|} \end{aligned}$$

- With $|L|$ the size of the lexicon (=set of unigrams).
- Why do we get $\epsilon|L|$ in the denominator and not just ϵ ?
- Recall that the number of occurrences of a $(n-1)$ -gram can be expressed as the number of occurrences of n -gram continuations.

$$\begin{aligned} C_{\text{smooth}}([w_{i-n+1}, \dots, w_{i-1}]) &= \sum_{w \in L} C_{\text{smooth}}([w_{i-n+1}, \dots, w_{i-1}, w]) \\ &= \sum_{w \in L} (C([w_{i-n+1}, \dots, w_{i-1}, w]) + \epsilon) \\ &= \sum_{w \in L} C([w_{i-n+1}, \dots, w_{i-1}, w]) + \epsilon|L| \\ &= C([w_{i-n+1}, \dots, w_{i-1}]) + \epsilon|L| \end{aligned}$$

Another solution: stupid backoff

- Instead of messing with all the n -gram counts, one can revert to a $(n - 1)$ -gram model, just in case the n -gram probability is 0.
- If for instance $P(\text{great}|\text{I had such a})$ is 0, then maybe $P(\text{great}|\text{had such a})$ is not? And if it is 0 too, then let's try $P(\text{great}|\text{such a})$! Etc.
- This kind of hybrid model can be expressed in the following, recursive³ way (λ being some discounting factor).

$$SB(w_i|[w_{i-n+1}, \dots, w_{i-1}]) = \begin{cases} P(w_i|[w_{i-n+1}, \dots, w_{i-1}]) & \text{if } > 0 \\ \lambda SB(w_i|[w_{i-n+2}, \dots, w_{i-1}]) & \text{otherwise} \end{cases}$$

- This model is no longer a proper probability distribution. Additionally, it requires us to compute k -gram occurrences for potentially every $k < n$!
- But unlike the smoothing solution, it may assign different probabilities to unlikely vs. super unlikely sequences.

³Notice how the “otherwise” case calls SB on a smaller context.

<https://colab.research.google.com/drive/1tCpAVF-SWnosVbRek8gpD9iz1g8lRhiNo>